

Question 7

A) Laplace  $(x|\mu, \sigma) = \frac{\sigma}{2} e^{-\sigma/2 |x-\mu|}$

negative log:  $\log\left(\frac{1}{\frac{\sigma}{2} e^{-\sigma/2 |x-\mu|}}\right)$   
 $(-\log(x) = \log(\frac{1}{x}))$

Differentiate

$$\frac{d}{dx} \left[ \log\left(\frac{1}{\frac{\sigma}{2} e^{-\sigma/2 |x-\mu|}}\right) \right]$$

$$= \frac{\sigma}{2 e^{\sigma/2 |x-\mu|}} \cdot \frac{d}{dx} \left[ \frac{2 e^{\sigma/2 |x-\mu|}}{\sigma} \right]$$

$$= \frac{\sigma}{2} \cdot \frac{d}{dx} [e^{\sigma/2 |x-\mu|}] \cdot \frac{2 e^{-\sigma/2 |x-\mu|}}{\sigma}$$

$$= e^{\sigma/2 |x-\mu|} \cdot \frac{d}{dx} [\sigma/2 |x-\mu|] \cdot e^{-\sigma/2 |x-\mu|}$$

$$= \sigma \cdot \frac{d}{dx} [|x-\mu|]$$

$$= \sigma \cdot \frac{x-\mu}{|x-\mu|} \cdot \frac{d}{dx} [x-\mu]$$

$$= \frac{\sigma \cdot \left( \frac{d}{dx} [x] + \frac{d}{dx} [-\mu] \right) (x-\mu)}{|x-\mu|}$$

$$= \frac{2 \cdot (1 + 0)(x - \mu)}{|x - \mu|} = \frac{2 \cdot (x - \mu)}{|x - \mu|}$$

$$|x - \mu|$$

$$|x - \mu|$$

1B)

$$\log [p(y | \vec{x}; \vec{w}) p(\vec{w})]$$

$$= \log [p(y | \vec{x}; \vec{w})] + \log (p(\vec{w}))$$

$$= \log [p(y | \vec{x}; \vec{w})] + \log \left[ \prod_{n=1}^K p(w_n) \right]$$

$$= \log [p(y | \vec{x}; \vec{w})] + \log \left[ \prod_{n=1}^K \frac{\sigma}{2} e^{-2|x|} \right]$$

$$= \log [p(y | \vec{x}; \vec{w})] + \sum_{n=1}^K \log \left( \frac{\sigma}{2} e^{-2|x|} \right)$$

$$7c) \max [p(y|\vec{x}; \vec{w}) p(\vec{w})]$$

$$= \max [\log(p(y|\vec{x}; \vec{w})) + \log(p(\vec{w}))]$$

This corresponds to  $L_2$  regularisation  
because  $p(\vec{w}) = \frac{\lambda}{2} e^{-\lambda \|\vec{w}\|^2}$  has an absolute value.

# Assignment 5

Sebastian Doka

2023-03-30

## Question 2

A)

```
coris = read.csv("coris.dat")
coris[["id"]] = NULL # remove "row.names" column
# Normalize the numeric columns:
normalise = function(x) {
  return ((x - mean(x)) / sd(x)) }
normcoris <- as.data.frame(lapply(coris, normalise))
# use the unnormalized versions of the binary variables:
normcoris$famhist <- coris$famhist
normcoris$chd <- coris$chd # this is the response variable
library(caret)

## Warning: package 'caret' was built under R version 4.1.3

## Loading required package: ggplot2

## Warning: package 'ggplot2' was built under R version 4.1.3

## Loading required package: lattice

set.seed(1)
ind <- createDataPartition(normcoris$chd, p = .6, list=FALSE)
train = normcoris[ind,]
test = normcoris[-ind,]

model <- glm(chd ~ ., data=train, family=binomial)
summary(model)

##
## Call:
## glm(formula = chd ~ ., family = binomial, data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
```

```
## -1.8537 -0.7747 -0.4293 0.8638 2.5509
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.15474    0.21026  -5.492 3.97e-08 ***
## sbp          0.03015    0.15495   0.195 0.84572
## tobacco      0.37981    0.15889   2.390 0.01683 *
## ldl          0.44514    0.16659   2.672 0.00754 **
## adiposity    0.15991    0.30610   0.522 0.60139
## famhist      0.75893    0.29746   2.551 0.01073 *
## typea        0.30349    0.15524   1.955 0.05058 .
## obesity     -0.30672    0.25298  -1.212 0.22535
## alcohol      0.13022    0.13808   0.943 0.34561
## age          0.72812    0.22941   3.174 0.00150 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 359.61  on 277  degrees of freedom
## Residual deviance: 281.52  on 268  degrees of freedom
## AIC: 301.52
##
## Number of Fisher Scoring iterations: 5
```

```
yprobs = predict(model, newdata=test, type='response')
ypred = round(yprobs)
table(ypred, test$chd)
```

```
##
## ypred  0  1
##      0 96 29
##      1 25 34
```

```
error_rate = mean(ypred != test$chd)
error_rate
```

```
## [1] 0.2934783
```

The variable with the most effect is the famhist because it has the highest estimated value(0.75893) meaning it has the most effect on the variable chd.

The miscalculation rate is: 0.2934783

```
library (randomForest)
```

```
## Warning: package 'randomForest' was built under R version 4.1.3
```

```
## randomForest 4.7-1.1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
## Attaching package: 'randomForest'

## The following object is masked from 'package:ggplot2':
##
##      margin

set.seed (1)
rf.coris <- randomForest (chd ~ ., data = train,mtry=6, importance = TRUE)

## Warning in randomForest.default(m, y, ...): The response has five or fewer
## unique values. Are you sure you want to do regression?

importance(rf.coris)

##              %IncMSE IncNodePurity
## sbp          2.1265043      6.418943
## tobacco     16.6329001     11.804149
## ldl          5.8404232      7.435982
## adiposity    6.2063156      5.943725
## famhist      5.5952444      2.331843
## typea       -0.6661421      4.838857
## obesity     -3.0476706      4.626176
## alcohol     -1.0952868      3.812399
## age         18.9628772     11.132201

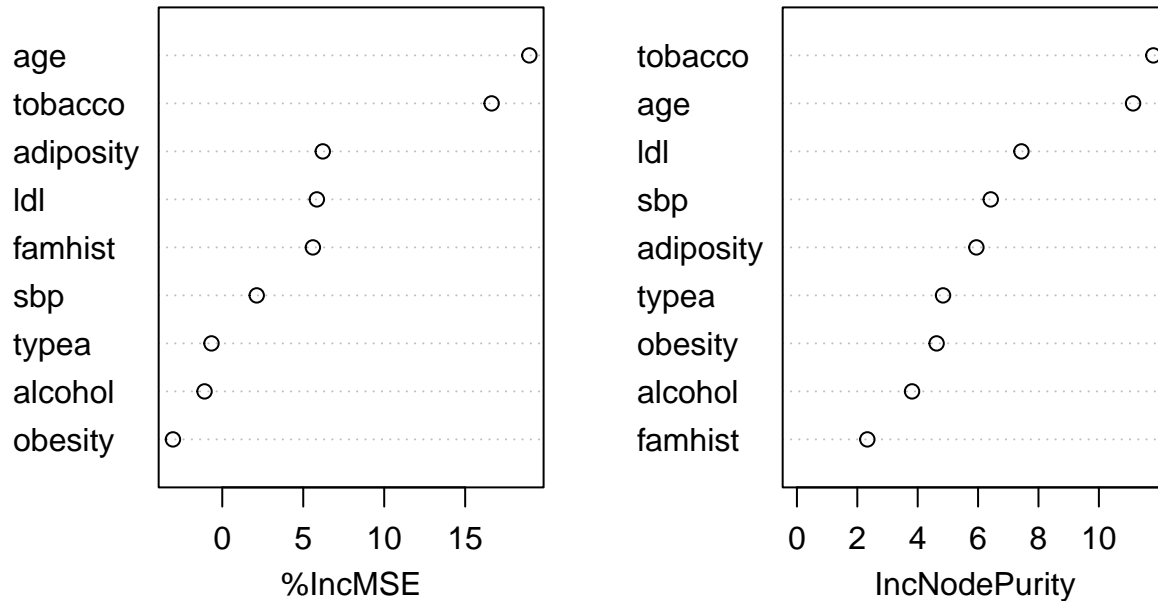
yprobs = predict(rf.coris, newdata=test, type='response')
ypred = round(yprobs)
miscalulation_rate = mean(ypred != test$chd)
print(paste("The miscalculation rate is: ",miscalulation_rate))

## [1] "The miscalculation rate is:  0.33695652173913"

varImpPlot(rf.coris)
```



## rf.coris



The variable with the greatest importance is tobacco with an importance of 16.6329001. The miscalculation rate is: 0.33695652173913.

Question3)


```
load('mnist23small.Rdata') # 1000 training digits, 1000 test
# When first loaded, the y column in each data frame is of type integer
# and takes values 1 or 2 (the digit shown in the image).
# Convert these integer variables into nominal variables, i.e. factors.
# This is needed by some classification algorithms
test$y = as.factor(test$y)
train$y = as.factor(train$y)
# "Normalise" the intensity of each pixel
normalise <- function(x) {
  return(x / 255)
}
train[,-1] = lapply(train[, -1], normalise)
test[,-1] = lapply(test[, -1], normalise)
# Note that "normalise" often means: centre (i.e. subtract the mean), and
# scale to variance 1, which is not what we've done here, but it's similar.

showdigit = function(imrow, label) {
  im = matrix(imrow, nrow=28)[,28:1] #reverse
  image(im, col=gray((0:255)/255),
  xaxt='n', yaxt='n', main=label, asp = 1
```

```

)
}
par(mfrow=c(2,4))
for(i in 1:8){
showdigit(as.matrix(train[i, 2:785]), paste(train$y[i]))
}

```



A)KNN with K=30

```
library(class)
```

```
## Warning: package 'class' was built under R version 4.1.3
```

```
library(caret)
```

```
Xtrain = train[,2:785]
```

```
Xtrain_matrix = data.matrix(Xtrain)
```

```
Xtest = test[,2:785]
```

```
Xtest_matrix = data.matrix(Xtest)
```

```
Ntest = 1000
```

```
ypred = knn(Xtrain_matrix , Xtest_matrix ,train$y, k = 30, prob = TRUE)
```

```
error.rate = function(ypred, ytrue){
```

```
err.rate = mean(ypred != ytrue)
return(err.rate)
}
print(paste('Error rate = ', error.rate(ypred,test$y)))
```

```
## [1] "Error rate = 0.026"
```

B)

```
library(randomForest)
set.seed(1)
bag.mnist <- randomForest(train$y ~ ., data = train, mtry = ncol(train)-1, importance = TRUE)

ypred <- predict (bag.mnist , newdata = test)

print(paste('Error rate = ', error.rate(ypred,test$y)))
```

```
## [1] "Error rate = 0.053"
```

C)

```
set.seed(1)
rf.mnist <- randomForest(train$y ~ ., data = train, mtry = sqrt(ncol(train)-1), importance = TRUE)

ypred <- predict (rf.mnist , newdata = test)

print(paste('Error rate = ', error.rate(ypred,test$y)))
```

```
## [1] "Error rate = 0.021"
```

D)

```
library(nnet)
mnist.nn2 <- nnet(train$y ~ ., data = train, size = 10, rang = 0.01,
decay = 0.01, maxit = 50, MaxNWts = 7862)
```

```
## # weights: 7861
## initial value 693.167695
## iter 10 value 50.936814
## iter 20 value 4.752770
## iter 30 value 2.954903
## iter 40 value 2.449441
## iter 50 value 2.335438
## final value 2.335438
## stopped after 50 iterations
```

```
ypred <- predict (mnist.nn2 , newdata = test, type="class")


print(paste('Error rate = ', error.rate(ypred,test$y)))
```


```
## [1] "Error rate = 0.036"
```

E)


```
for (i in 1:10){  
  w=mnist.nn2$wts  
  recep = w[(785*(i-1)+1):(785*(i-1)+784)]  
  showdigit(as.matrix(recep), paste(i))}
```

1







3



4




5





6



7

